



Welcome

CRUD with NodeJS and Python and Datastax Astra





The Crew

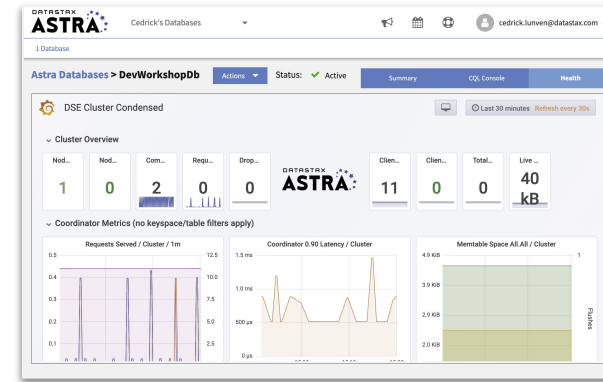


DataStax Developer Advocacy Special Unit


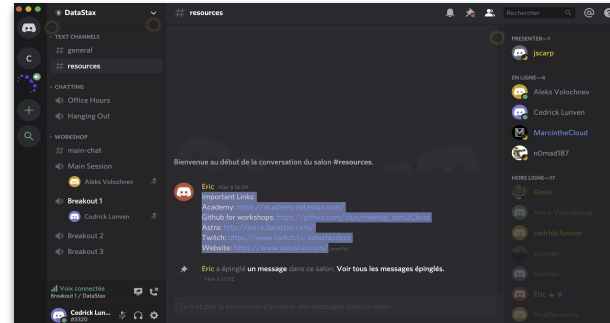
Courses: youtube.com/DataStaxDevs




Runtime: dtsx.io/workshop



Questions: bit.ly/cassandra-workshop



Quizz: menti.com



Disclaimer

This is a coding session. You will need some experience with the NodeJS or Python and have a github account.

You can do these exercises on a local install, but we will only use Gitpod today during the workshop.

No need to install anything.



Hands-on exercise material



Get your instance here:

- <http://dtsx.io/workshop>

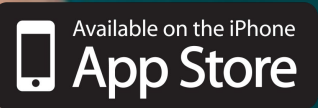
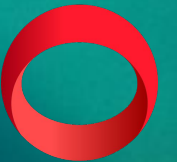


Repository:

- <https://github.com/DataStax-Academy/workshop-crud-with-python-and-node>



menti.com



Application Development **CRUD**



**What we will
cover:**





1. Use Case and data model
2. Set up Astra database and schema
3. Connect to Astra
4. Create and update records
5. Read results

Showtime !!



Journey Checklist

- ☒ Writing 1000 Rows to Astra
- ☒ Reading 1000 Rows from Astra
- ☐ Replaying Rows

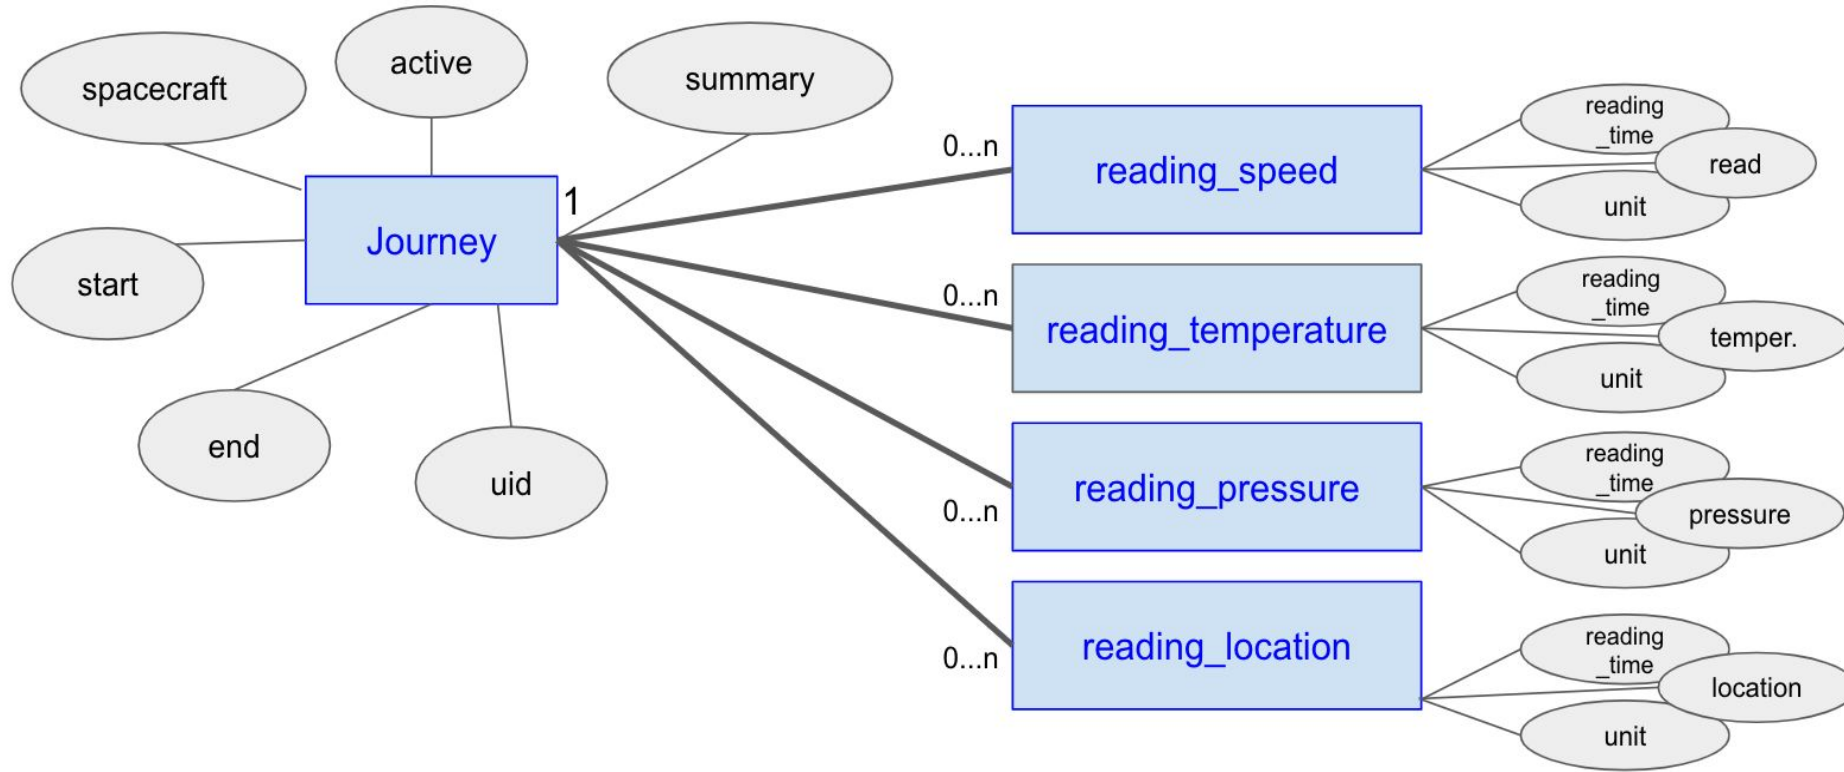

Astra Table	Rows Written	Rows Read	Current Row Displayed
spacecraft_temperature_over_time	1000	1000	189
spacecraft_speed_over_time	1000	1000	189
spacecraft_pressure_over_time	1000	1000	189
spacecraft_location_over_time	1000	1000	189



Designing your data model

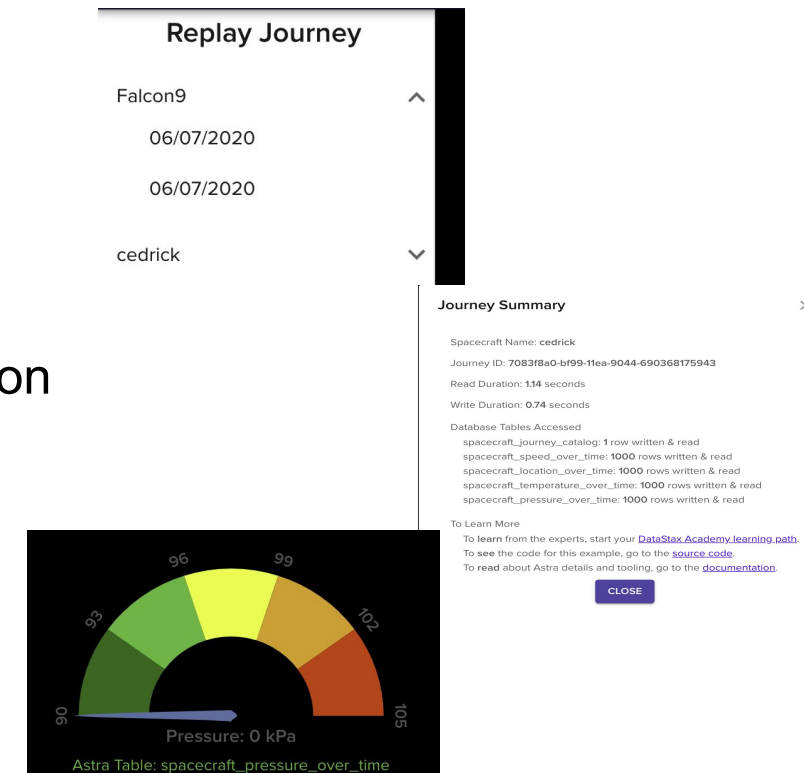



#1 Conceptual Data Model




#1 Application Workflow

- **Space crafts catalog queries**
 - Look up all of the journeys for a spacecraft
 - Look up the state of a journey
 - Create a new journey
- **Sensor readings queries : Speed, Pressure, Temperature, Location**
 - Save readings over time
 - Analyze each dimension independently
 - Analyze data per journey
 - Less than 100.000 records per journey per dimension



#2 Map to Logical Data Model



spacecraft_journey_catalog


spacecraft_temperature_over_time

spacecraft_location_over_time

spacecraft_speed_over_time

spacecraft_pressure_over_time

#3 Logical Data Model



spacecraft_journey_catalog

spacecraft_name	K
journey_id	C ↓
start	
end	
active	
summary	

spacecraft_temperature_over_time

spacecraft_speed_over_time


spacecraft_name	K
journey_id	K
reading_time	C ↓
speed	
speed_unit	

spacecraft_location_over_time

spacecraft_pressure_over_time

spacecraft_name	K
journey_id	K
reading_time	C ↓
pressure	
pressure_unit	

#4 Physical Data Model



spacecraft_journey_catalog

spacecraft_name	text
journey_id	timeuuid
start	timestamp
end	timestamp
active	boolean
summary	text

spacecraft_temperature_over_time

spacecraft_speed_over_time


spacecraft_name	text
journey_id	timeuuid
reading_time	timestamp
speed	double
speed_unit	text

spacecraft_location_over_time

spacecraft_pressure_over_time

spacecraft_name	text
journey_id	timeuuid
reading_time	timestamp
pressure	double
pressure_unit	text


#5 CQL DDL



```
CREATE TABLE IF NOT EXISTS spacecraft_journey_catalog (  
  spacecraft_name text,  
  journey_id      timeuuid,  
  start           timestamp,  
  end             timestamp,  
  active          boolean,  
  summary         text,  
  PRIMARY KEY ((spacecraft_name),  
  WITH CLUSTERING ORDER BY (journey_id ASC));
```

```
CREATE TABLE IF NOT EXISTS spacecraft_speed_over_time (  
  spacecraft_name text,  
  journey_id      timeuuid,  
  speed           double,  
  reading_time    timestamp,  
  speed_unit      text,  
  PRIMARY KEY ((spacecraft_name, journey_id), reading_time))  
  WITH CLUSTERING ORDER BY (reading_time DESC);
```

#5 CQL DDL with UDT



```
CREATE TYPE IF NOT EXISTS location_udt (  
    x_coordinate double,  
    y_coordinate double,  
    z_coordinate double  
);
```

```
CREATE TABLE IF NOT EXISTS spacecraft_location_over_time (  
    spacecraft_name text,  
    journey_id timeuuid,  
    reading_time timestamp,  
    location frozen<location_udt>,  
    location_unit text,  
    PRIMARY KEY ((spacecraft_name, journey_id), reading_time)  
);
```


Application Development **CRUD**



**What we will
cover:**

1. Use Case and Datamodel
2. Set up Astra database and schema
3. Connect to Astra
4. Create and update records
5. Read results

Introduction to Astra



==



Global Scale

Put your data where you need it without compromising performance, availability, or accessibility.



No Operations

Eliminate the overhead to install, operate, and scale Cassandra.



5 Gig Free Tier

Launch a database in the cloud with a few clicks, no credit card required.

Hands-on prep work: Create your Astra database

- If you are new to Astra, sign up here: **dtsx.io/workshop**
- add a new keyspace (if you already have one)
- Take a note of the username, password, keyspace of your cluster
- We will be using these defaults:
 - keyspace: **spacecraft**
 - user name: **SUser**
 - password: **SPassword1**
- Download the secure-connect-bundle, **this is unique to you and your database instance**

Hands-on Exercise 1:

- **Create the Astra database**
- **Create the schema in the database**
 - **Log into Astra CQL console**
 - **Copy schema from github and paste into CQL console**

Application Development **CRUD**



**What we will
cover:**

1. Use Case and Datamodel
2. Set up Astra database and schema
3. **Connect to Astra**
4. Create and update records
5. Read results

Datastax Drivers

One of set drivers to connect them all - January 2020



Connectivity

- Token & Datacenter Aware
- Load Balancing Policies
- Retry Policies
- Reconnection Policies
- Connection Pooling
- Health Checks
- Authentication | Authorization
- SSL

Query

- CQL Support
- Schema Management
- Sync/Async/Reactive API
- Query Builder
- Compression
- Paging

Parsing Results

- Lazy Load
- Object Mapper
- Spring Support
- Paging

Install Drivers

```
npm install cassandra-driver
```



```
pip install cassandra-driver
```



Connection to Cassandra ...with ASTRA

```
const client = new cassandra.Client({  
  cloud: { secureConnectBundle: 'secure.zip' },  
  credentials: { username: 'u', password: 'p' }  
});
```



```
auth_provider = PlainTextAuthProvider(  
    username='U', password='P')  
cluster = Cluster(  
    Cloud = { Secure_connect_bundle: 'secure.zip'},  
    auth_provider=auth_provider, protocol_version=2)  
session= cluster.connect('killrvideo')
```



Important about the session

- It is a stateful object handling communications with each node
- session should be unique per application (Singleton)
- session should be closed at application shutdown (shutdown hook) in order to free opened TCP sockets (stateful)

Node: `client.shutdown();`

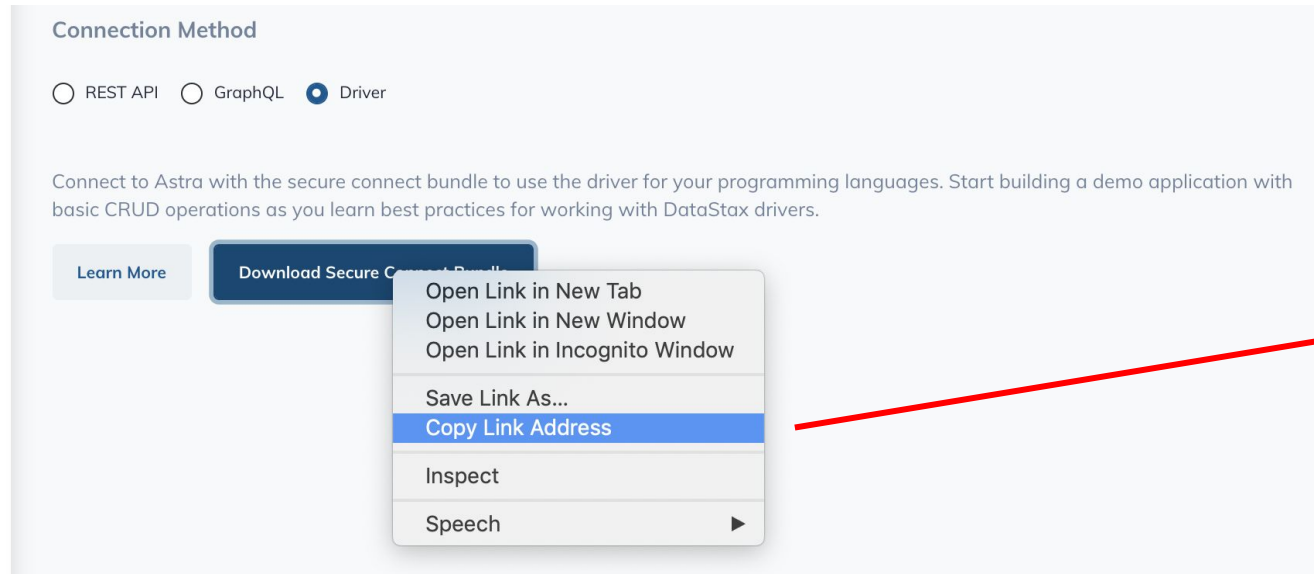
Python: `session.shutdown();`

Hands-on Exercise 2:

Setup connection to Astra

- **Get gitpod ready**
- **Install the drivers**
- **Save the secure connection bundle in gitpod**
- **Update the connection settings in node and python**
- **Test the connection:**
 - **Run Ex02 code**

Get the secure bundle with curl



click to copy the link to
clipboard
(refresh page before copy,
link is only 'fresh' for short
time)

```
curl -L "<insert link here>" > creds.zip
```

Application Development **CRUD**



**What we will
cover:**

1. Use Case and Datamodel
2. Set up Astra database and schema
3. Connect to Astra
4. **Create and update records**
5. Read results

Execute simple statements

```
client.execute('select * from t1 where c1 = ?', [5]);
```



```
session.execute("select * from t1 where c1 = %s", 5);
```



Prepared and Bound Statements

- Compiled once on each node automatically as needed
- Prepare each statement only once per application

```
const query = 'select * from t1 where c1 =?;';  
client.execute(query, [5], { prepare: true })
```



```
prepared = session.prepare("select * from t1 where c1=?")  
result = session.execute(prepared, [5])
```



Hands-on Exercise 3 and 4:

Inserts with simple and prepared statements

- **Simple insertion into spacecraft_journey_catalog:**
 - **Run Ex03 code**
- **Insertion with prepared statement:**
 - **Adapt Ex04 code and run**

Working with batch statements

```
var myBatch = [  
  { query: insertSpeed, params: [spacecraft_name, journey_id, speed, readingTime, 'km/hour' ] },  
  { query: insertTemperature, params: [spacecraft_name, journey_id, pressure, readingTime, 'Pa' ] },  
  { query: insertPressure, params: [spacecraft_name, journey_id, temperature, readingTime, 'K' ] },  
  { query: insertLocation, params: [spacecraft_name, journey_id, location, readingTime, 'AU' ] }  
]  
  
const result = await connection.client.batch(myBatch, { prepare: true })
```



```
batch = BatchStatement()  
batch.add(prepared_insertLocation, [spacecraft_name, journey_id, Location(x,y,z), readingTime, 'AU' ])  
batch.add(prepared_insertSpeed, [spacecraft_name, journey_id, speed, readingTime, 'km/hour' ])  
batch.add(prepared_insertTemperature, [spacecraft_name, journey_id, pressure, readingTime, 'Pa' ])  
batch.add(prepared_insertPressure, [spacecraft_name, journey_id, temperature, readingTime, 'K' ])  
connection.session.execute(batch)
```



Working with UDTs (Python special)

- Recommended to register UDTs with Cluster instance
- When using prepared statements, not necessary to register

```
class Foo(object):  
    def __init__(self, street, zipcode, otherstuff):  
        self.street = street  
        self.zipcode = zipcode  
        self.otherstuff = otherstuff  
  
insert_statement = session.prepare("INSERT INTO users (id, location) VALUES (?, ?)")  
  
session.execute(insert_statement, [0, Foo("123 Main St.", 78723, "some other stuff")])
```



Hands-on Exercise 5:

Inserts with UDTs and batches

- **Adapt and run Ex05**

Application Development **CRUD**



**What we will
cover:**



1. Use Case and Datamodel
2. Set up Astra database and schema
3. Connect to Astra
4. Create and update records
5. **Read results**

Hands-on Exercise 7:

Simple selects, parsing results and paging

- **Simple selects: Ex07**
- **Parsing: Ex08 and Ex09**
- **Paging: Ex10**

menti.com




Engage !

Share with us you Cassandra use cases !

Share with your vision and future of Cloud Native

Share what you need to succeed with Cassandra.



Developer Resources

LEARN

- Join academy.datastax.com
- Browse www.datastax.com/dev

ASK/SHARE

Join community.datastax.com

Ask/answer community user questions - share your expertise

CONNECT

Follow us

We are on Youtube - Twitter - Twitch!

MATERIALS

Slides and materials from all workshops are available at <https://github.com/DataStax-Academy>



Thank You

